



PRESENTED BY



strataconf.com

#StrataHadoop

Modeling big data with R, sparklyr, and Apache Spark

1:30pm–5:00pm Tuesday, March 14, 2017

Data science & advanced analytics

Location: LL21 C/D

Level: Intermediate

Secondary topics: R

John Mount (Win-Vector LLC)

Steve Nolen (RStudio)

Edgar Ruiz (RStudio)

url: <https://github.com/WinVector/BigDataRStrata2017>

RStudio Shiny Server Pro Accounts

- Distribute credentials and get everyone started with RStudio Server Pro.
- Server time generously donated by RStudio.



RStudio Server Pro

RStudio for the Enterprise

RStudio is the premier IDE for R. RStudio Server lets you access RStudio from anywhere using a web browser. RStudio Server Pro delivers the team productivity, security, centralized management, metrics, and commercial support that professional data science teams need to develop at scale.

From: <https://www.rstudio.com/products/rstudio-server-pro/>

What are we going to do?

- Become `dplyr` masters.

Work through markdowns in a bit

- Exercises/01-Universal-tools.Rmd
- Exercises/02-Big-Data.Rmd
- slides/lazyeval.Rmd

If you have trouble: ask your neighbors, flag me and the TAs, or peek in Exercises/solutions. These exercises are a memory aid, not a test.

Exercises/01-Universal-tools.Rmd

Exercise 1


- Let's start with Exercises/01-Universal-tools.Rmd
- ~~We will just read through “Exercise 1” a bit.~~

Exercises/01-Universal-tools.Rmd

Exercise 2

iris

On average, which species has the greatest difference in *petal width* and *petal length*?



Exercises/01-Universal-tools.Rmd

Exercise 2

1. **Group** iris by Species
2. For each group, return:
 - `avg_width = mean Petal.Width`
 - `avg_length = mean Petal.Length`
3. For each row, calculate `diff = avg_length - avg_width`
4. Return the **row** whose `diff == the max diff`
5. Return the **columns** above

Hint 1: save each results in a different variable such as `iris1`, `iris2`, ... and examine your results before moving to the next step.

Hint 2: Consider save the max diff *value* into a temp table before trying to find which row has that value.

05:00

Exercises/01-Universal-tools.Rmd

Exercise 2 solution

```
library(dplyr)
```

```
iris1 <- group_by(iris, Species)
```

```
iris2 <- summarise(iris1,  
  avg_width = mean(Petal.Width),  
  avg_length = mean(Petal.Length))
```

```
iris3 <- mutate(iris2, diff = avg_length - avg_width)
```

```
iris4 <- filter(iris3, diff == max(diff))
```

```
select(iris4, Species, avg_width, avg_length)
```

Mini-topic pipes and pipelines

% ≥ %

Ceci n'est pas une pipe.

The pipe operator

%>%

```
filter(iris, Sepal.Length == max(Sepal.Length))
```

```
iris %>% filter(., Sepal.Length == max(Sepal.Length))
```

```
iris %>% filter(Sepal.Length == max(Sepal.Length))
```

These all do the
same thing
Try it!



Exercises/01-Universal-tools.Rmd

Exercise 3

Use `%>%` to turn your code from that last exercise into a single long pipe.

03:00

Exercises/01-Universal-tools.Rmd

Exercise 3 solution:

Take Exercise 2 solution

```
library(dplyr)
```

```
iris1 <- group_by(iris, Species)
```

```
iris2 <- summarise(iris1,  
  avg_width = mean(Petal.Width),  
  avg_length = mean(Petal.Length))
```

```
iris3 <- mutate(iris2, diff = avg_length - avg_width)
```

```
iris4 <- filter(iris3, diff == max(diff))
```

```
select(iris4, Species, avg_width, avg_length)
```

Exercises/01-Universal-tools.Rmd

Exercise 3 solution:

and *mechanically* translate it into a pipeline

```
iris %>%  
  group_by(Species) %>%  
  summarise(  
    avg_width = mean(Petal.Width),  
    avg_length = mean(Petal.Length)) %>%  
  mutate(diff = avg_length - avg_width) %>%  
  filter(diff == max(diff)) %>%  
  select(Species, avg_width, avg_length)
```

pipeline debugging hint

- Break the pipeline early and look at intermediate results.
- Use “->” (“write arrow”) to save result to “.” (dot).
 - Since right arrow is forbidden by most style guides you can search for it to make sure you have not left debugging code in!

Debugging example

```
iris %>%
```

```
  group_by(Species) %>%
```

```
  summarise(
```

```
    avg_width = mean(Petal.Width),
```

```
    avg_length = mean(Petal.Length)) ->.
```

```
print(.)
```

```
. %>% mutate(diff = avg_length - avg_width) %>%
```

```
filter(diff == max(diff)) %>%
```

```
select(Species, avg_width, avg_length)
```



~~Exercises/01-Universal-tools.Rmd~~

~~Exercise 4~~

- ~~• Try joins on the “band data.”~~

03:00

Exercises/01-Universal-tools.Rmd


Exercise 5

nycflights13



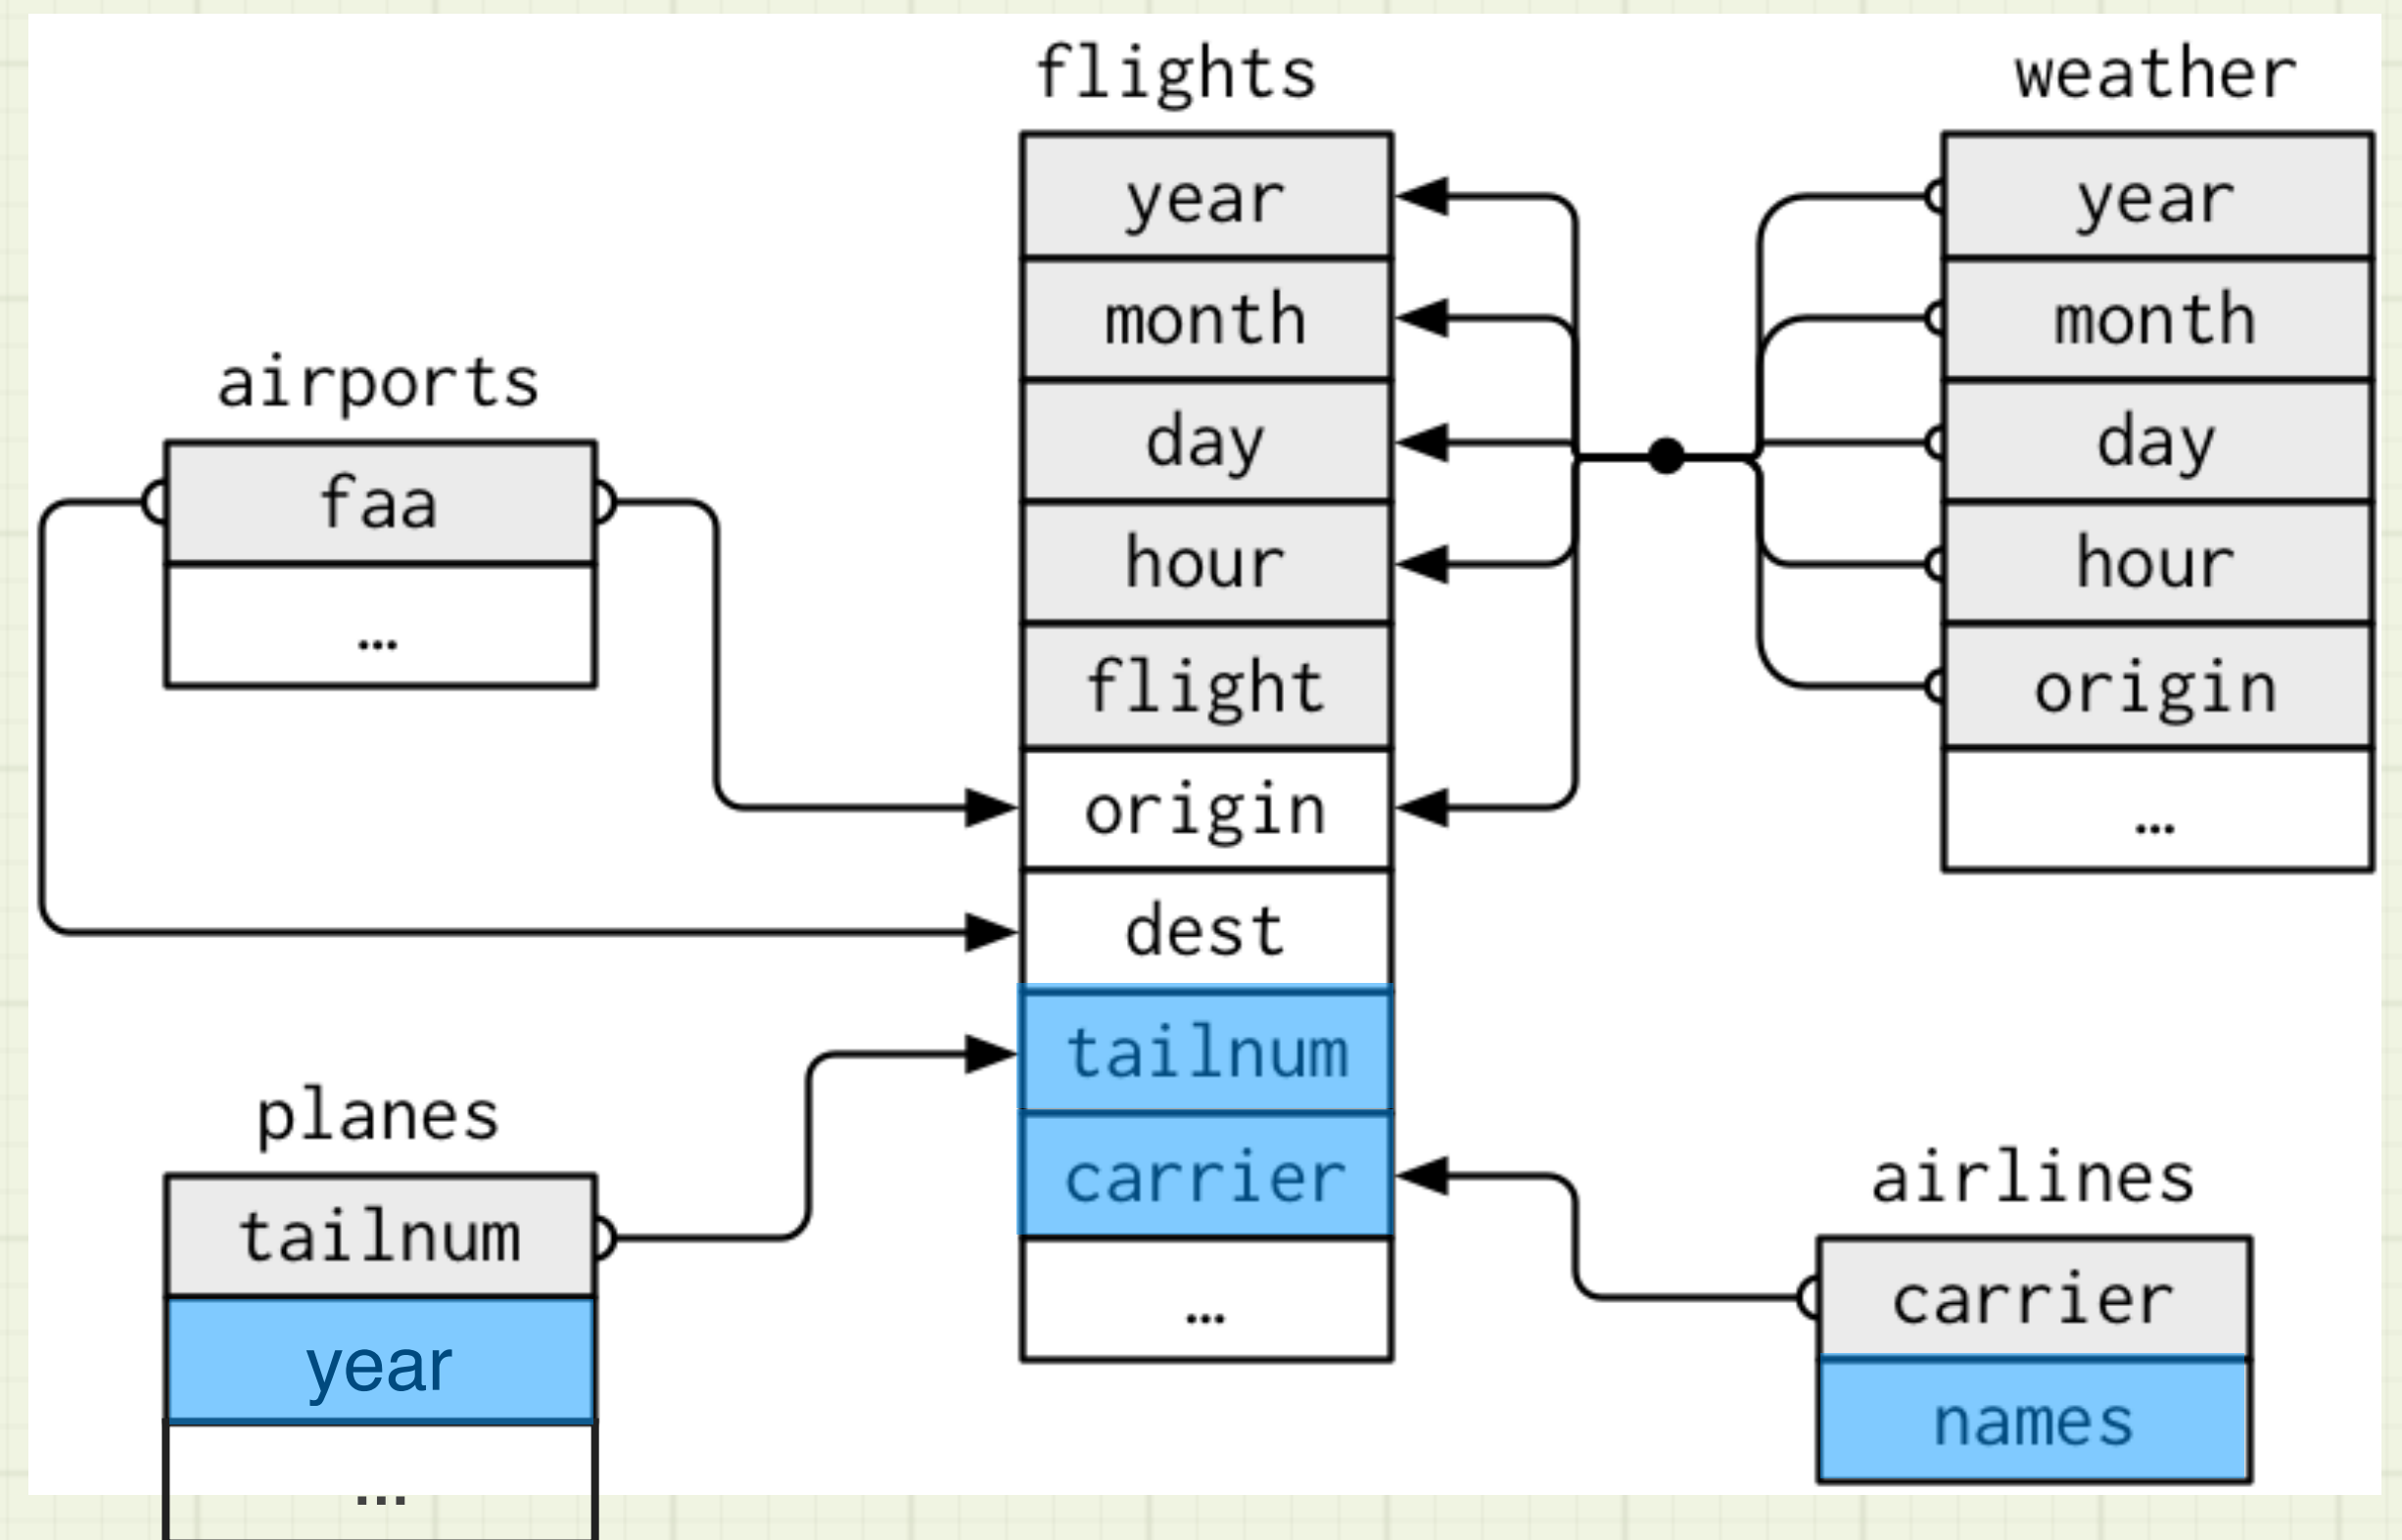
Data on every flight that departed La Guardia, JFK, or Newark airports in 2013

nycflights13



nycflights13

On average, which airline has the newest planes (assigned to the NYC area)?



Exercises/01-Universal-tools.Rmd

Exercise 5

Determine which airline has the newest planes. Please start with the code in the notebook.

```
flights %>%  
  distinct(carrier, tailnum) %>%  
  _____ %>%  
  _____ %>%  
  _____ ...
```

10:00

Exercises/01-Universal-tools.Rmd

Exercise 5 answer

name	avg	n	nas
Hawaiian Airlines Inc.	2011.77	14	1
Virgin America	2008.71	53	1
Frontier Airlines Inc.	2008.00	26	3
Alaska Airlines Inc.	2007.84	84	1
JetBlue Airways	2006.50	193	6
SkyWest Airlines Inc.	2005.86	28	0
Endeavor Air Inc.	2004.71	204	2
Mesa Airlines Inc.	2003.56	58	1
ExpressJet Airlines Inc.	2002.44	316	8
AirTran Airways Corporation	2002.21	129	17

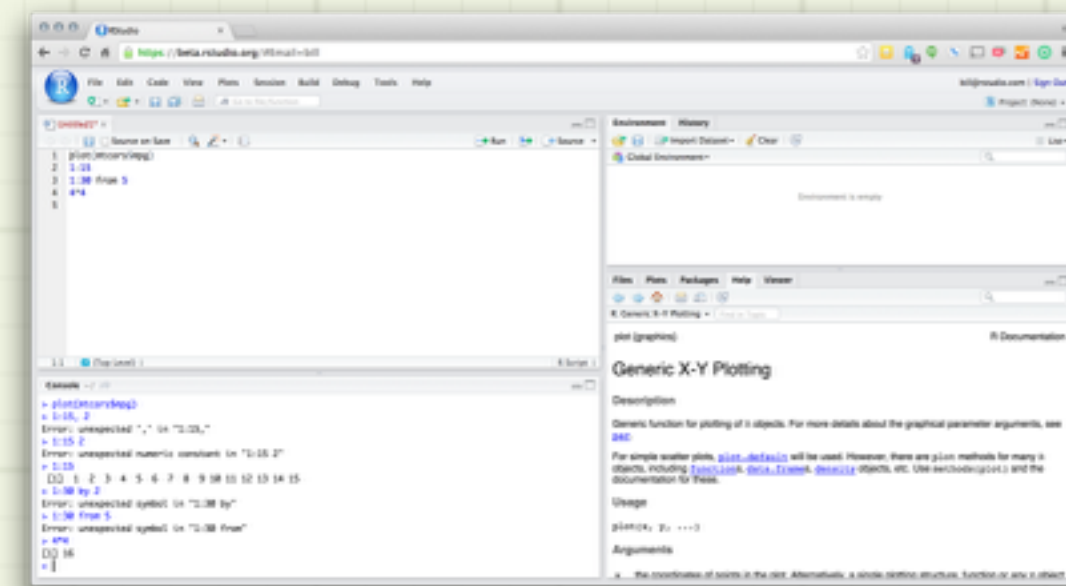
Exercises/01-Universal-tools.Rmd

Exercise 5 solution

```
flights %>%  
  # selects distinct combinations of carrier and tailnum  
  distinct(carrier, tailnum) %>%  
  # join to planes to get year manufactured  
  # (which column should you join on?)  
  left_join(planes, by = "tailnum") %>%  
  # group by carrier (e.g. the airline)  
  group_by(carrier) %>%  
  # calculate by carrier:  
  #   1. avg - the mean year (with na.rm = TRUE)  
  #   2. n - the total number of planes  
  #   3. nas - the number of planes with unknown year (year == NA)  
  summarise(avg = mean(year, na.rm = TRUE),  
             n = n(), nas = sum(is.na(year))) %>%  
  # join to airlines to get full airline name  
  # (which column should you join on?)  
  left_join(airlines, by = "carrier") %>%  
  # select just the name, avg, n, and nas variables in that order  
  select(name, avg, n, nas) %>%  
  # order the results by avg with the newest planes at the top  
  arrange(desc(avg))
```


Mini topic databases

User Browser



Server

RStudio Server
Pro



Database

Redshift



 Airlines

dplyr driver functions

Package	DBMS
<code>src_sqlite()</code>	SQLite
<code>src_mysql()</code>	MySQL, MariaDB
<code>src_postgres()</code>	PostgreSQL
<code>library(bigrquery)</code> <code>src_bigquery()</code>	Google BigQuery

<https://cran.r-project.org/web/packages/dplyr/vignettes/databases.html>

dplyr adapts to databases

- Many common base-R task must be translated dplyr to work on databases.
- Most dplyr verbs work with many data sources.

Base R	dplyr
<code>x\$col</code>	<code>select(x, col)</code>



dplyr database workflow

1. Create a connection

```
con <- dplyr::src_sqlite(":memory:", create = TRUE)
```

2. Create a reference

```
tab <- copy_to(con, data, 'tablename')
```

3. Manipulate the reference

```
query <- tab %>% filter(x > 1) %>% select(x, y, z)
```

4. Collect the results

```
results <- collect(query)
```

5. Close the connection

```
rm(con); gc()
```


1. Create a connection

```
con <- dplyr::src_sqlite(":memory:", create = TRUE)
```

Save
to use

src_driver
function

driver
specific args

Lists tables
in DB

```
src_tbls(con)
```

```
## "iris"      "iris2"      "iris3"
```

2. Create a table reference

connection
to DB

name of
table in DB

```
tab <- tbl(con, "table_name")
```

Use `tbl()` to create objects that refer to tables in the database

3. Manipulate the reference

Treat the reference as if it were a table in R. dplyr will translate your code to SQL and execute it in the DBMS.*

```
flights <- tbl(air, "flights")  
flights %>%  
  distinct(uniquecarrier, tailnum) %>%  
  ...
```

4. Collect the results

Use `collect()` to import the entire set of results into R.

```
q6 <- flights %>%  
  filter(year > 2007, depdelay > 15) %>%  
  filter(depdelay == 240) %>%  
  collect()
```

5. Close the connection

```
rm(air)  
gc()
```

dplyr automatically closes connections when you remove the connection object *and then run the garbage collector, gc().*

Exercises/02-Big-Data.Rmd

Determine which airline has the newest planes using `dplyr` to control data in our practice database.

```
flights %>%  
  distinct(carrier, tailnum) %>%  
  _____ %>%  
  _____ %>%  
  _____ ...
```

Hint: **CHEAT!!!**

03:00

Target start: 2:32pm

Mini-topic: lazy evaluation



“Just in Time” delivery of results.

slides/lazyeval.Rmd



Lazy Execution 1

```
q1 <- filter(flights, year < 2007)
q2 <- filter(q1, depdelay > 15)
q3 <- filter(q2, depdelay < 240)
q4 <- select(q3, arrdelay, depdelay, year)
q4
```

Lazy Execution 1

```
q1 <- filter(flights, year < 2007)
q2 <- filter(q1, depdelay > 15)
q3 <- filter(q2, depdelay < 240)
q4 <- select(q3, arrdelay, depdelay, year)
q4
```

dplyr will not retrieve data until last possible moment. It combines all necessary work into a single optimized query.

```
show_query(q4)
```

```
## <SQL>
```

```
## SELECT "arrdelay" AS "arrdelay",  
        "depdelay" AS "depdelay",  
        "year" AS "year"
```

```
## FROM "flights"
```

```
## WHERE "year" > 2007.0  
        AND "depdelay" > 15.0  
        AND "depdelay" < 240.0
```


collapse()

Forces execution in **DBMS**

```
q5 <- flights %>%  
  mutate(adjdelay = depdelay - 15) %>%  
  collapse() %>%  
  filter(adjdelay > 0)
```

collapse() turns the preceding queries into a table expression

remaining queries are run against the table described in the collapsed expression

Next:
Spark and sparklyr