



PRESENTED BY



strataconf.com

#StrataHadoop

Modeling big data with R, sparklyr, and Apache Spark

1:30pm–5:00pm Tuesday, March 14, 2017

Data science & advanced analytics

Location: LL21 C/D

Level: Intermediate

Secondary topics: R

John Mount (Win-Vector LLC)

Steve Nolen (RStudio)

Edgar Ruiz (RStudio)

url: <https://github.com/WinVector/BigDataRStrata2017>

3 levels of “big data”

1. Need to build summaries or simple predictive models from a large amount of data.

- Can sample data and run in memory.

2. Need to apply a model or transformation to a lot of data.

- Can distribute work to a cluster.

3. Need to compute something involving all the relations between the data.

- Need big data systems and algorithms.

Not teaching OPS/ENG

- We are loaning you single cluster accounts that are non-durable
 - temp tables disappear when you disconnect
- This means we (hopefully) don't end up running “a consulting clinic for installation bugs” (thanks Garrett Grolemond).
- It also means we are not discussing issues such as lifetime of data inside Spark and h2o, starting and stopping the cluster, pinning items into memory, or organization of storage.
- It isn't that OPS isn't important, it is just too big to be in this workshop (and varies depending on where you end up working with big data).



Configuration conflict



Big data systems have a cost

- Big data systems usually trade high throughput for some combination of:
 - High latency
 - Limited state
 - Limited communication patterns
 - Inconvenience
- In fact they can prohibit some superior algorithms



“Scalability! But at what COST?”

McSherry, Isard, Murray

- Define COST as cluster size (in machines) needed to be as fast as a good in-ram (single machine, multiple CPU) implementation.
- Typical COST of “best of breed” systems in article:
 - 10
 - 100
 - 256
 - infinity

Redefine “big data”

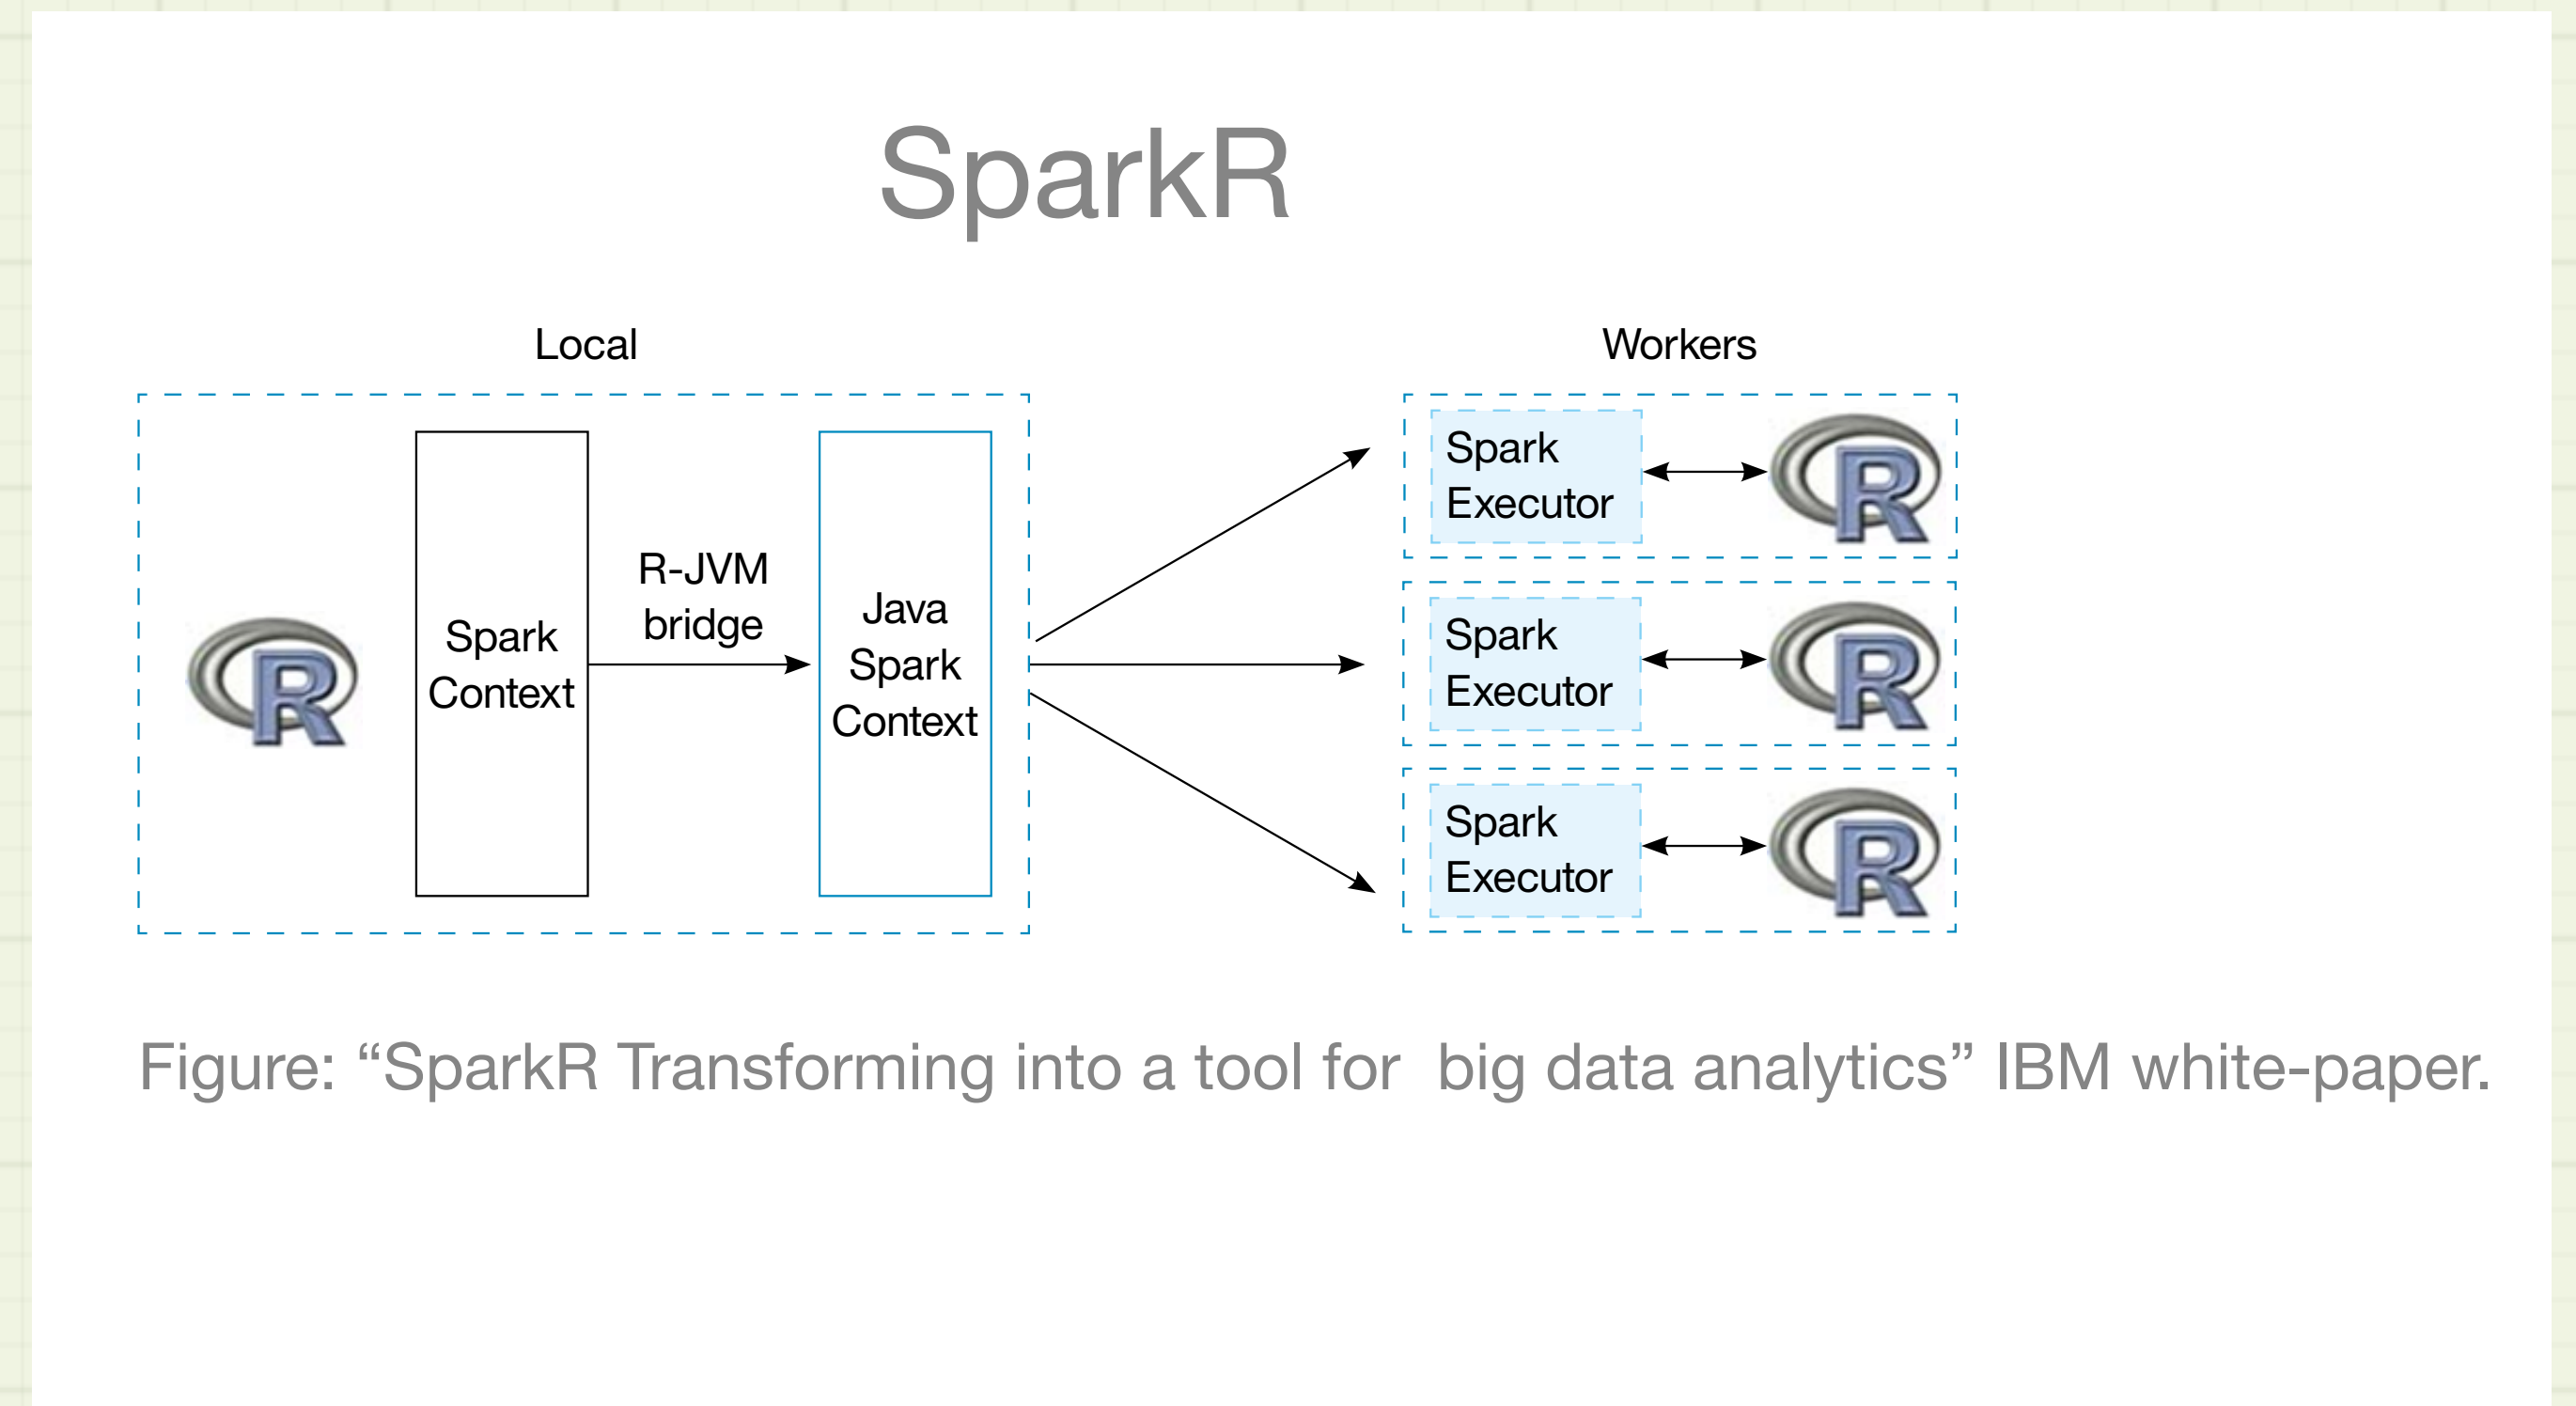

Big Data > 1/3 RAM

What is Spark?

- Open-source Apache computing engine
- Bigger-than-memory data, low-latency distributed computing
- Can integrate with the Hadoop ecosystem
- Built-in machine learning



Architecture choices

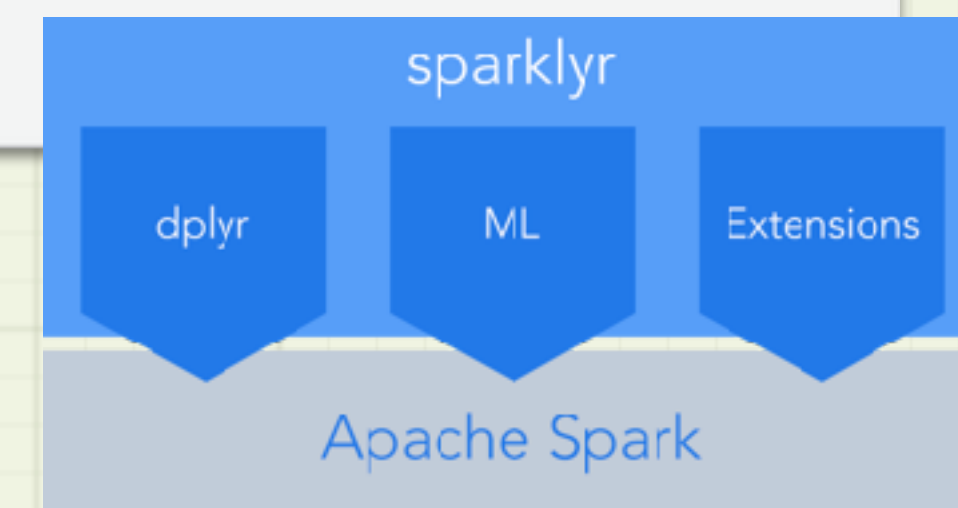
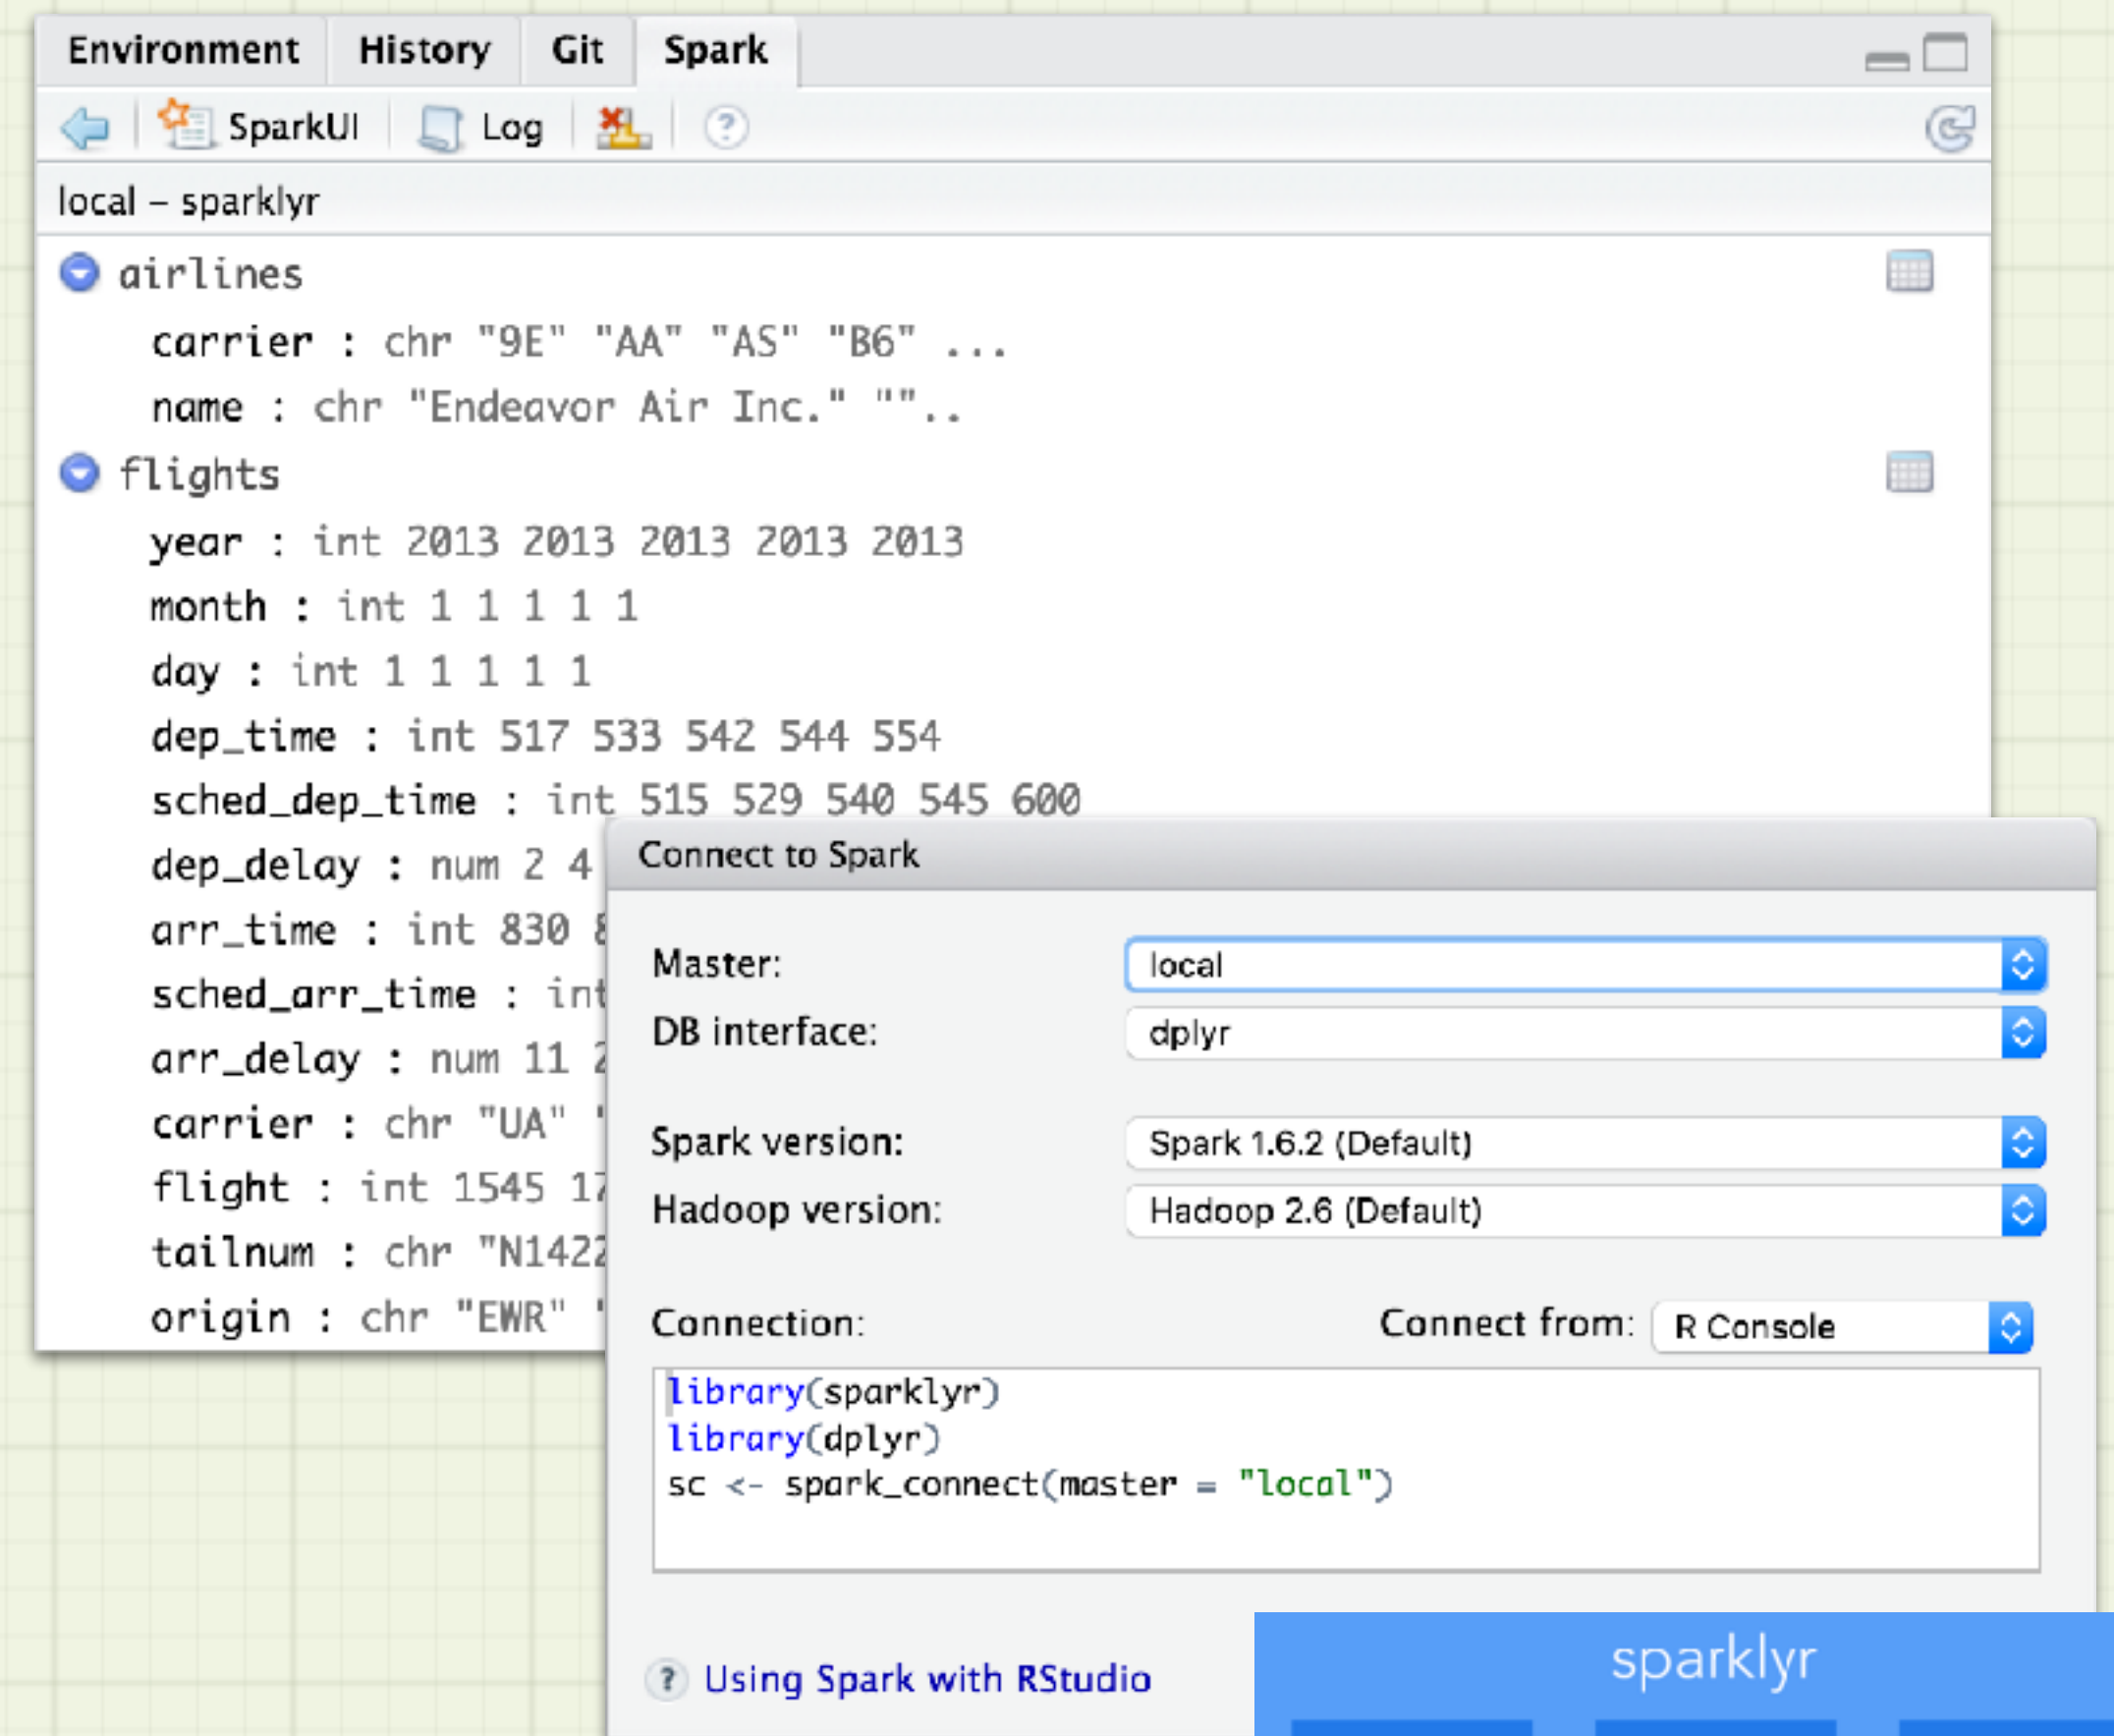


Comparison

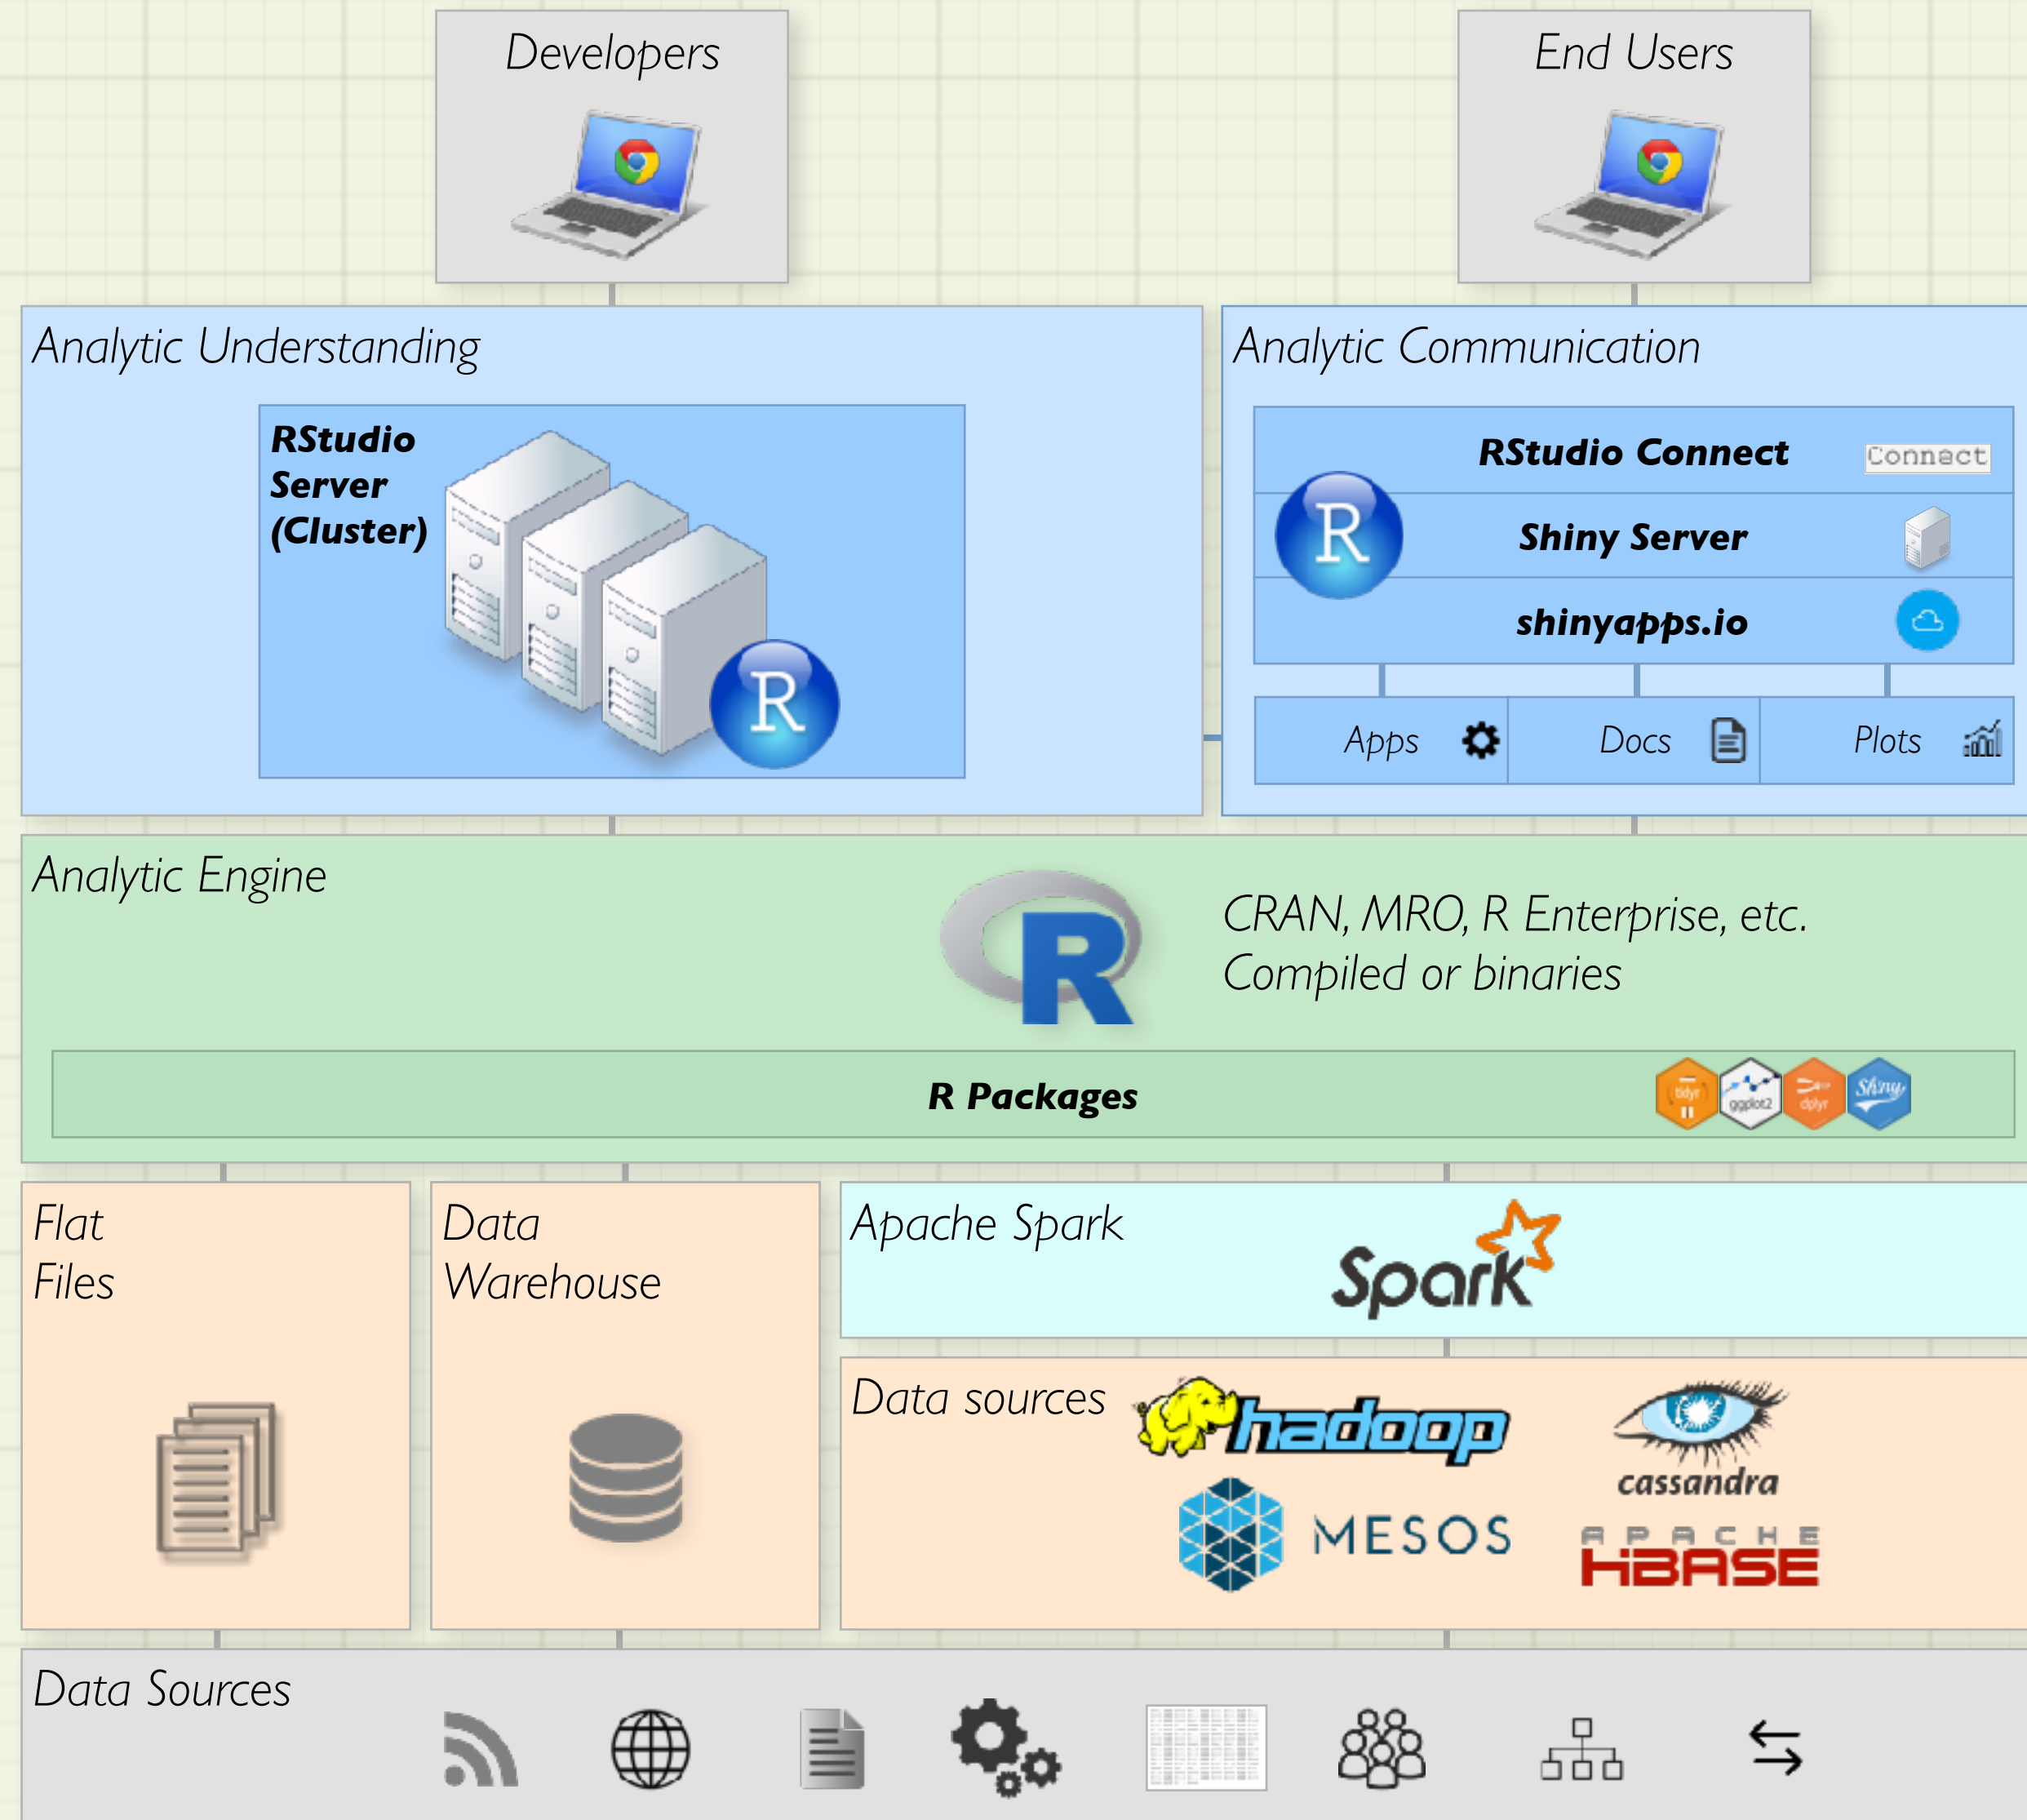
- SparklyR / h2o
 - Command remote native Spark using adapted `dplyr` notation to allow workflows *as if* you were working in R.
 - Use native SparkML, and h2o methods to do the work.
- SparkR
 - Use R to command copies of distributed R across a cluster.
 - Allows R user defined functions.
 - Not currently a `dplyr` backend.
- Which is “better”?
 - It depends on your data, infrastructure, tasks, and legacy code.
 - Also depends *a lot* on the current state of each adapter, and both are under rapid development.
 - Expect to see both going forward.
- We are going to concentrate on SparklyR.

sparklyr

- **R package.** New, open-source package from RStudio
- **dplyr.** Complete dplyr back-end for Spark
- **IDE.** Integrated with the RStudio IDE
- **Extensible.** Extensible foundation for Spark and R



Typical setup



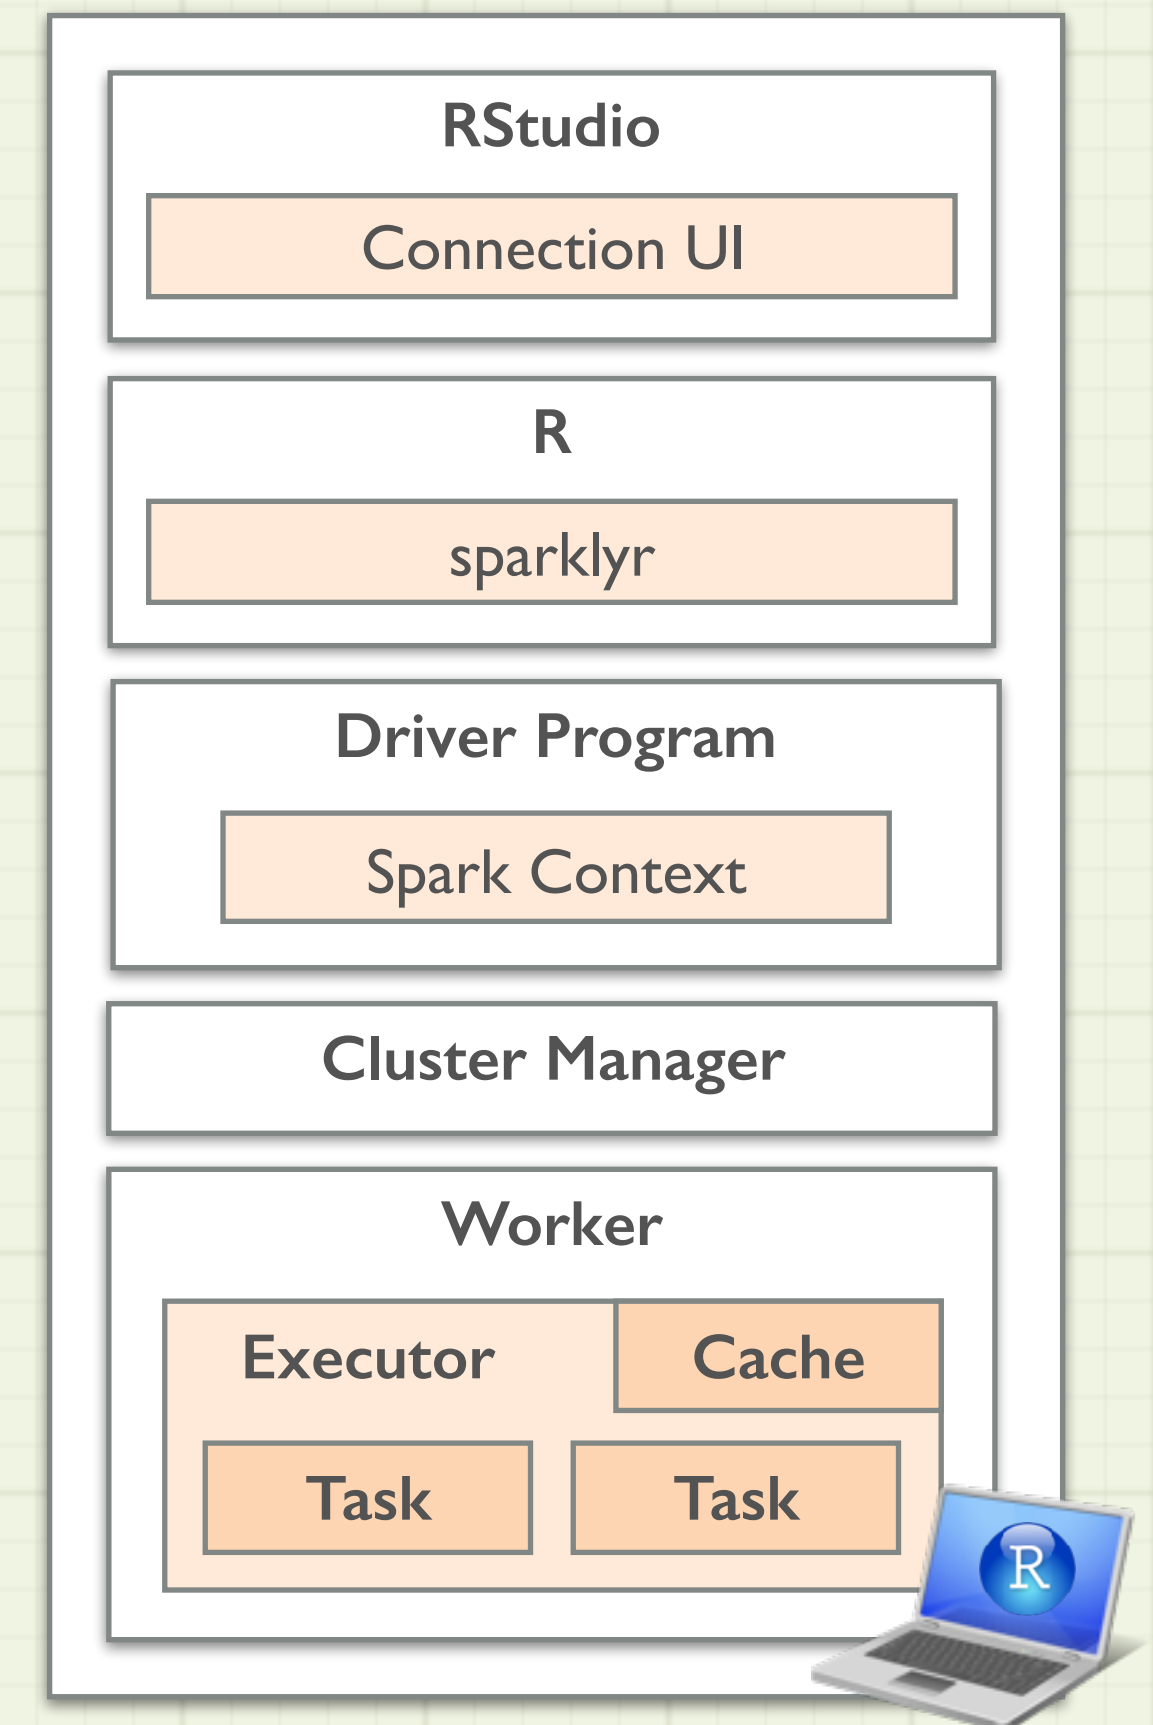
Local Mode

```
library(sparklyr)
library(dplyr)

spark_install()

sc <- spark_connect("local")

my_tbl <- copy_to(sc, iris)
```



Spark SQL

Use dplyr syntax to translate R code into Spark SQL (HiveQL)

DPLYR

```
my_tbl %>%  
  filter(Petal_Width < 0.3) %>%  
  select(Petal_Length, Petal_Width) %>%  
  arrange(Petal_Length)
```

SPARK SQL

```
select Petal_Length, Petal_Width  
from iris  
where Petal_Width < 0.3  
order by Petal_Length
```

Nota bene

- Remote data services (intentionally) look a lot like local data.
- But there are differences.



Examples

```
library("sparklyr")  
library("dplyr")  
sc <- spark_connect(master = "local",  
                     version = "2.0.0",  
                     hadoop_version="2.7")  
iris_tbl <- copy_to(sc, iris, "iris",  
                   overwrite= TRUE)
```

iris_tbl is a handle, not data

```
> str(iris)
```

```
'data.frame': 150 obs. of  5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

```
> str(iris_tbl)
```

```
List of 2
 $ src:List of 1
  ..$ con:List of 11
   .. ..$ master      : chr "local[4]"
   .. ..$ method      : chr "shell"
   .. ..$ app_name     : chr "sparklyr"
   .. ..$ config       :List of 5
   .. .. ..$ sparklyr.cores.local : int 4
 ...
```



fix: dplyr::glimpse()

```
> glimpse(iris_tbl)
```

```
Observations: 150
```

```
Variables: 5
```

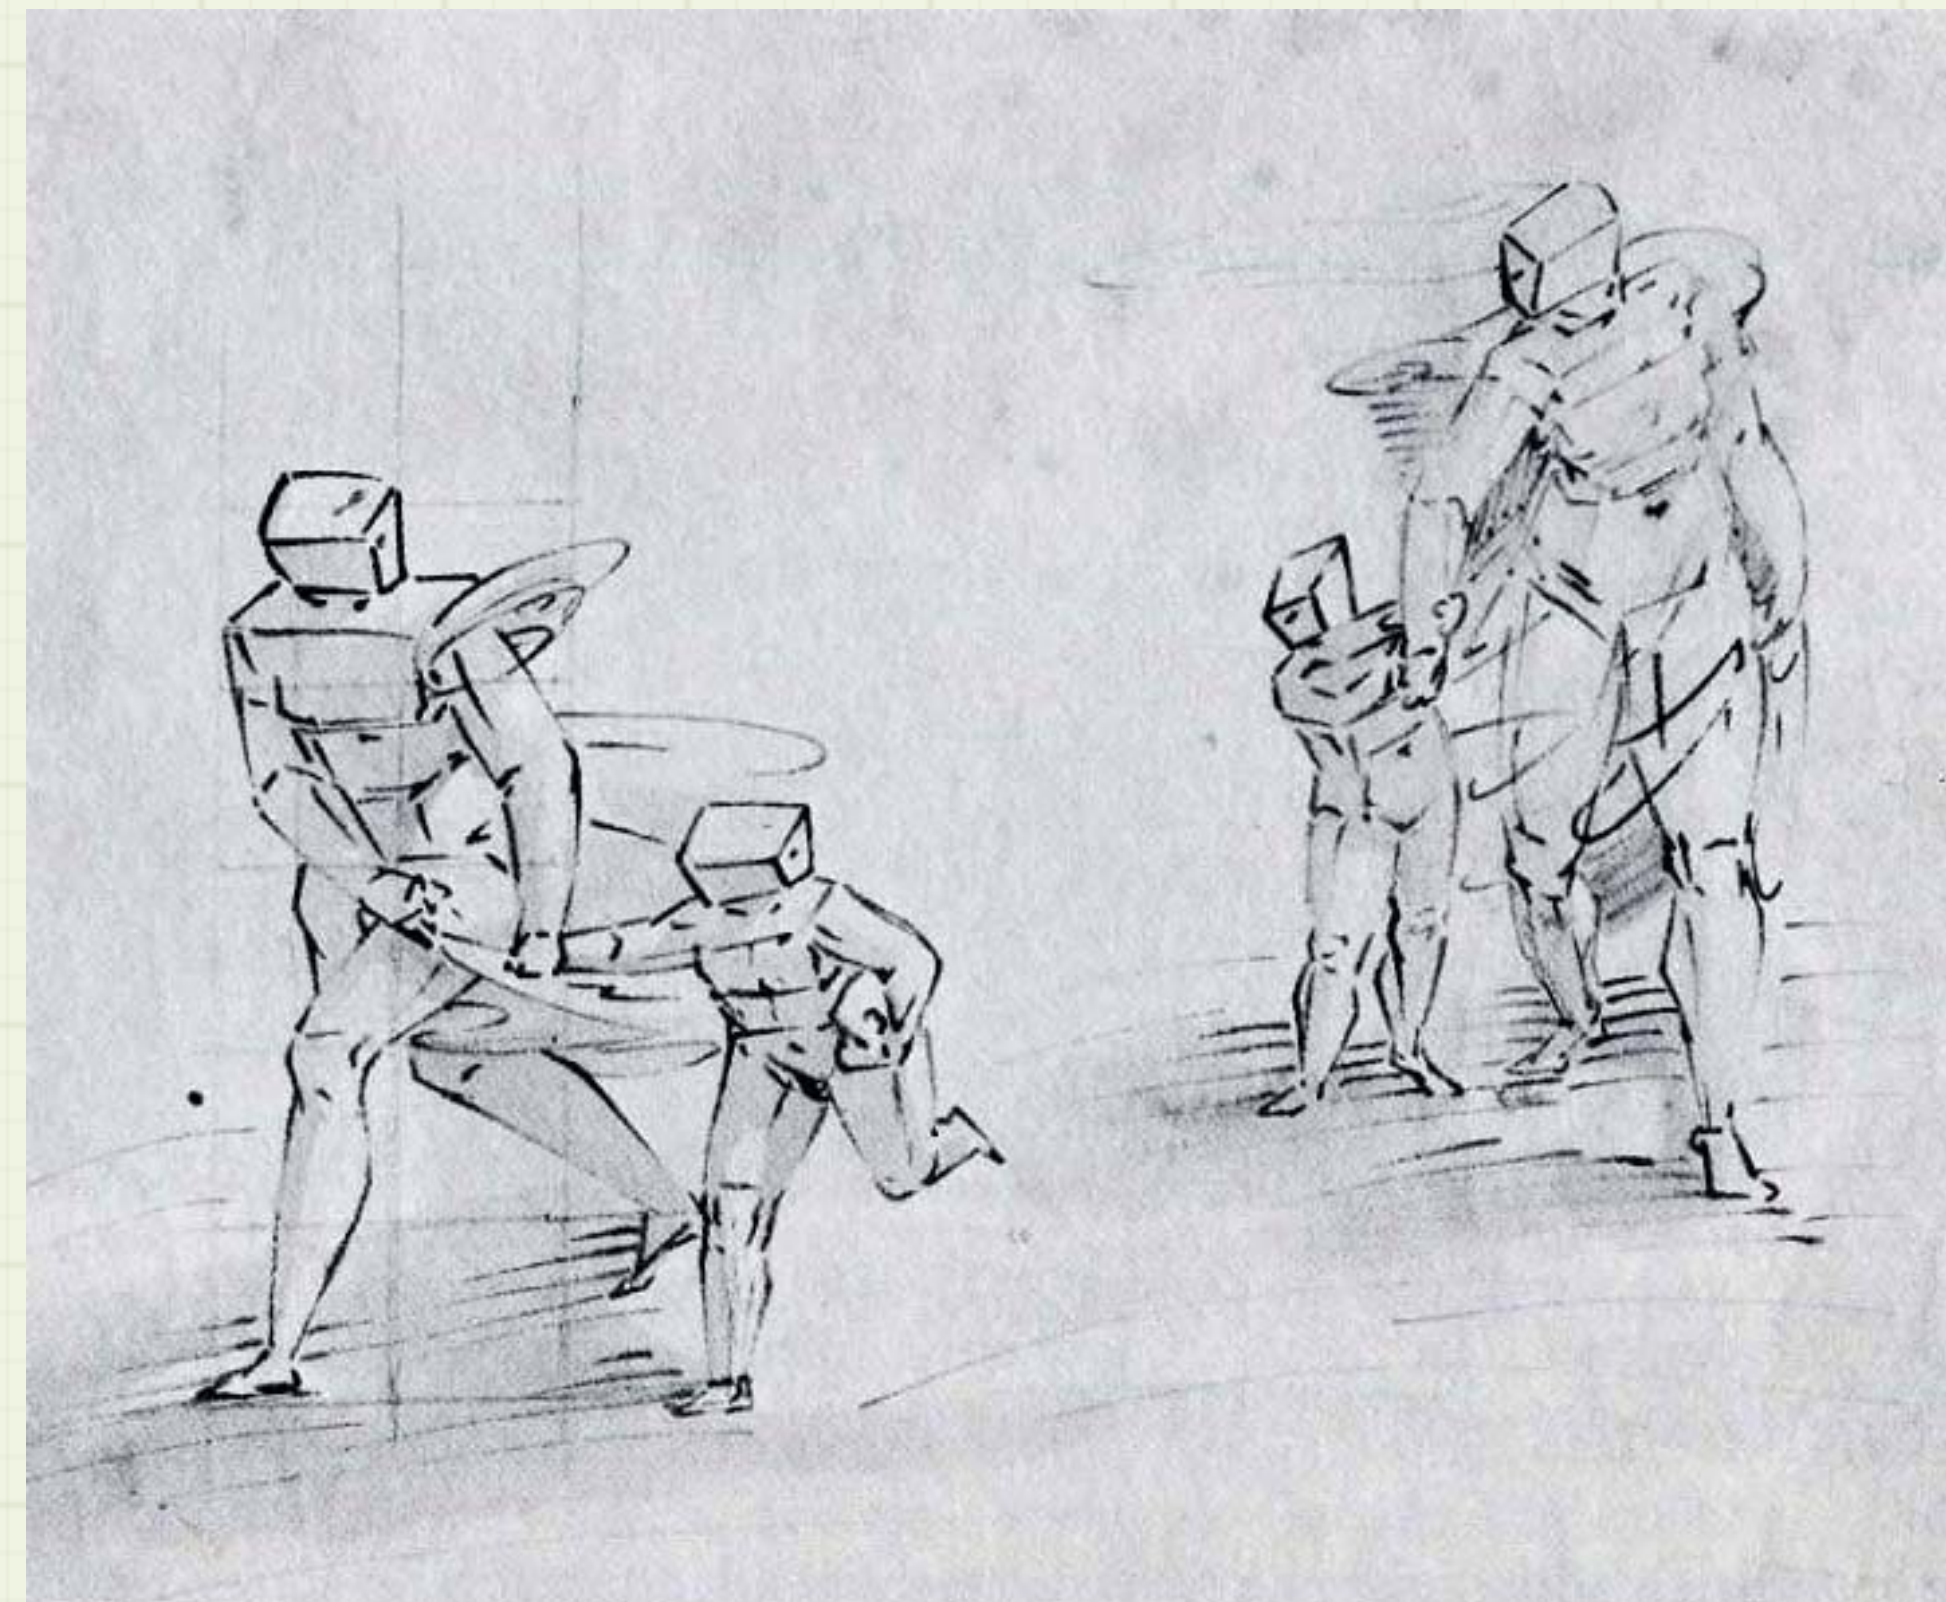
```
$ Sepal_Length <dbl> 5.1, 4.9, 4.7, 4.6, 5.0, 5.4, 4.6, 5.0, 4.4, 4.9, 5.4, 4.8, 4.8, 4.3, 5.8, 5.7, ...  
$ Sepal_Width  <dbl> 3.5, 3.0, 3.2, 3.1, 3.6, 3.9, 3.4, 3.4, 2.9, 3.1, 3.7, 3.4, 3.0, 3.0, 4.0, 4.4, ...  
$ Petal_Length <dbl> 1.4, 1.4, 1.3, 1.5, 1.4, 1.7, 1.4, 1.5, 1.4, 1.5, 1.5, 1.6, 1.4, 1.1, 1.2, 1.5, ...  
$ Petal_Width  <dbl> 0.2, 0.2, 0.2, 0.2, 0.2, 0.4, 0.3, 0.2, 0.2, 0.1, 0.2, 0.2, 0.1, 0.1, 0.2, 0.4, ...  
$ Species      <chr> "setosa", "setosa", "setosa", "setosa", "setosa", "setosa", "setosa", "setosa", ...
```


Mnemonic

In standard R operations are wrestling in with your data



With handles you send operations out to your data



Handles (def: abstract reference to a resource) necessarily have reference semantics. Though we can make objects in the data pool express value semantics.

again: *handles*

```
> summary(iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100	setosa :50
1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300	versicolor:50
Median :5.800	Median :3.000	Median :4.350	Median :1.300	virginica :50
Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199	
3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800	
Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500	

```
> summary(iris_tbl)
```

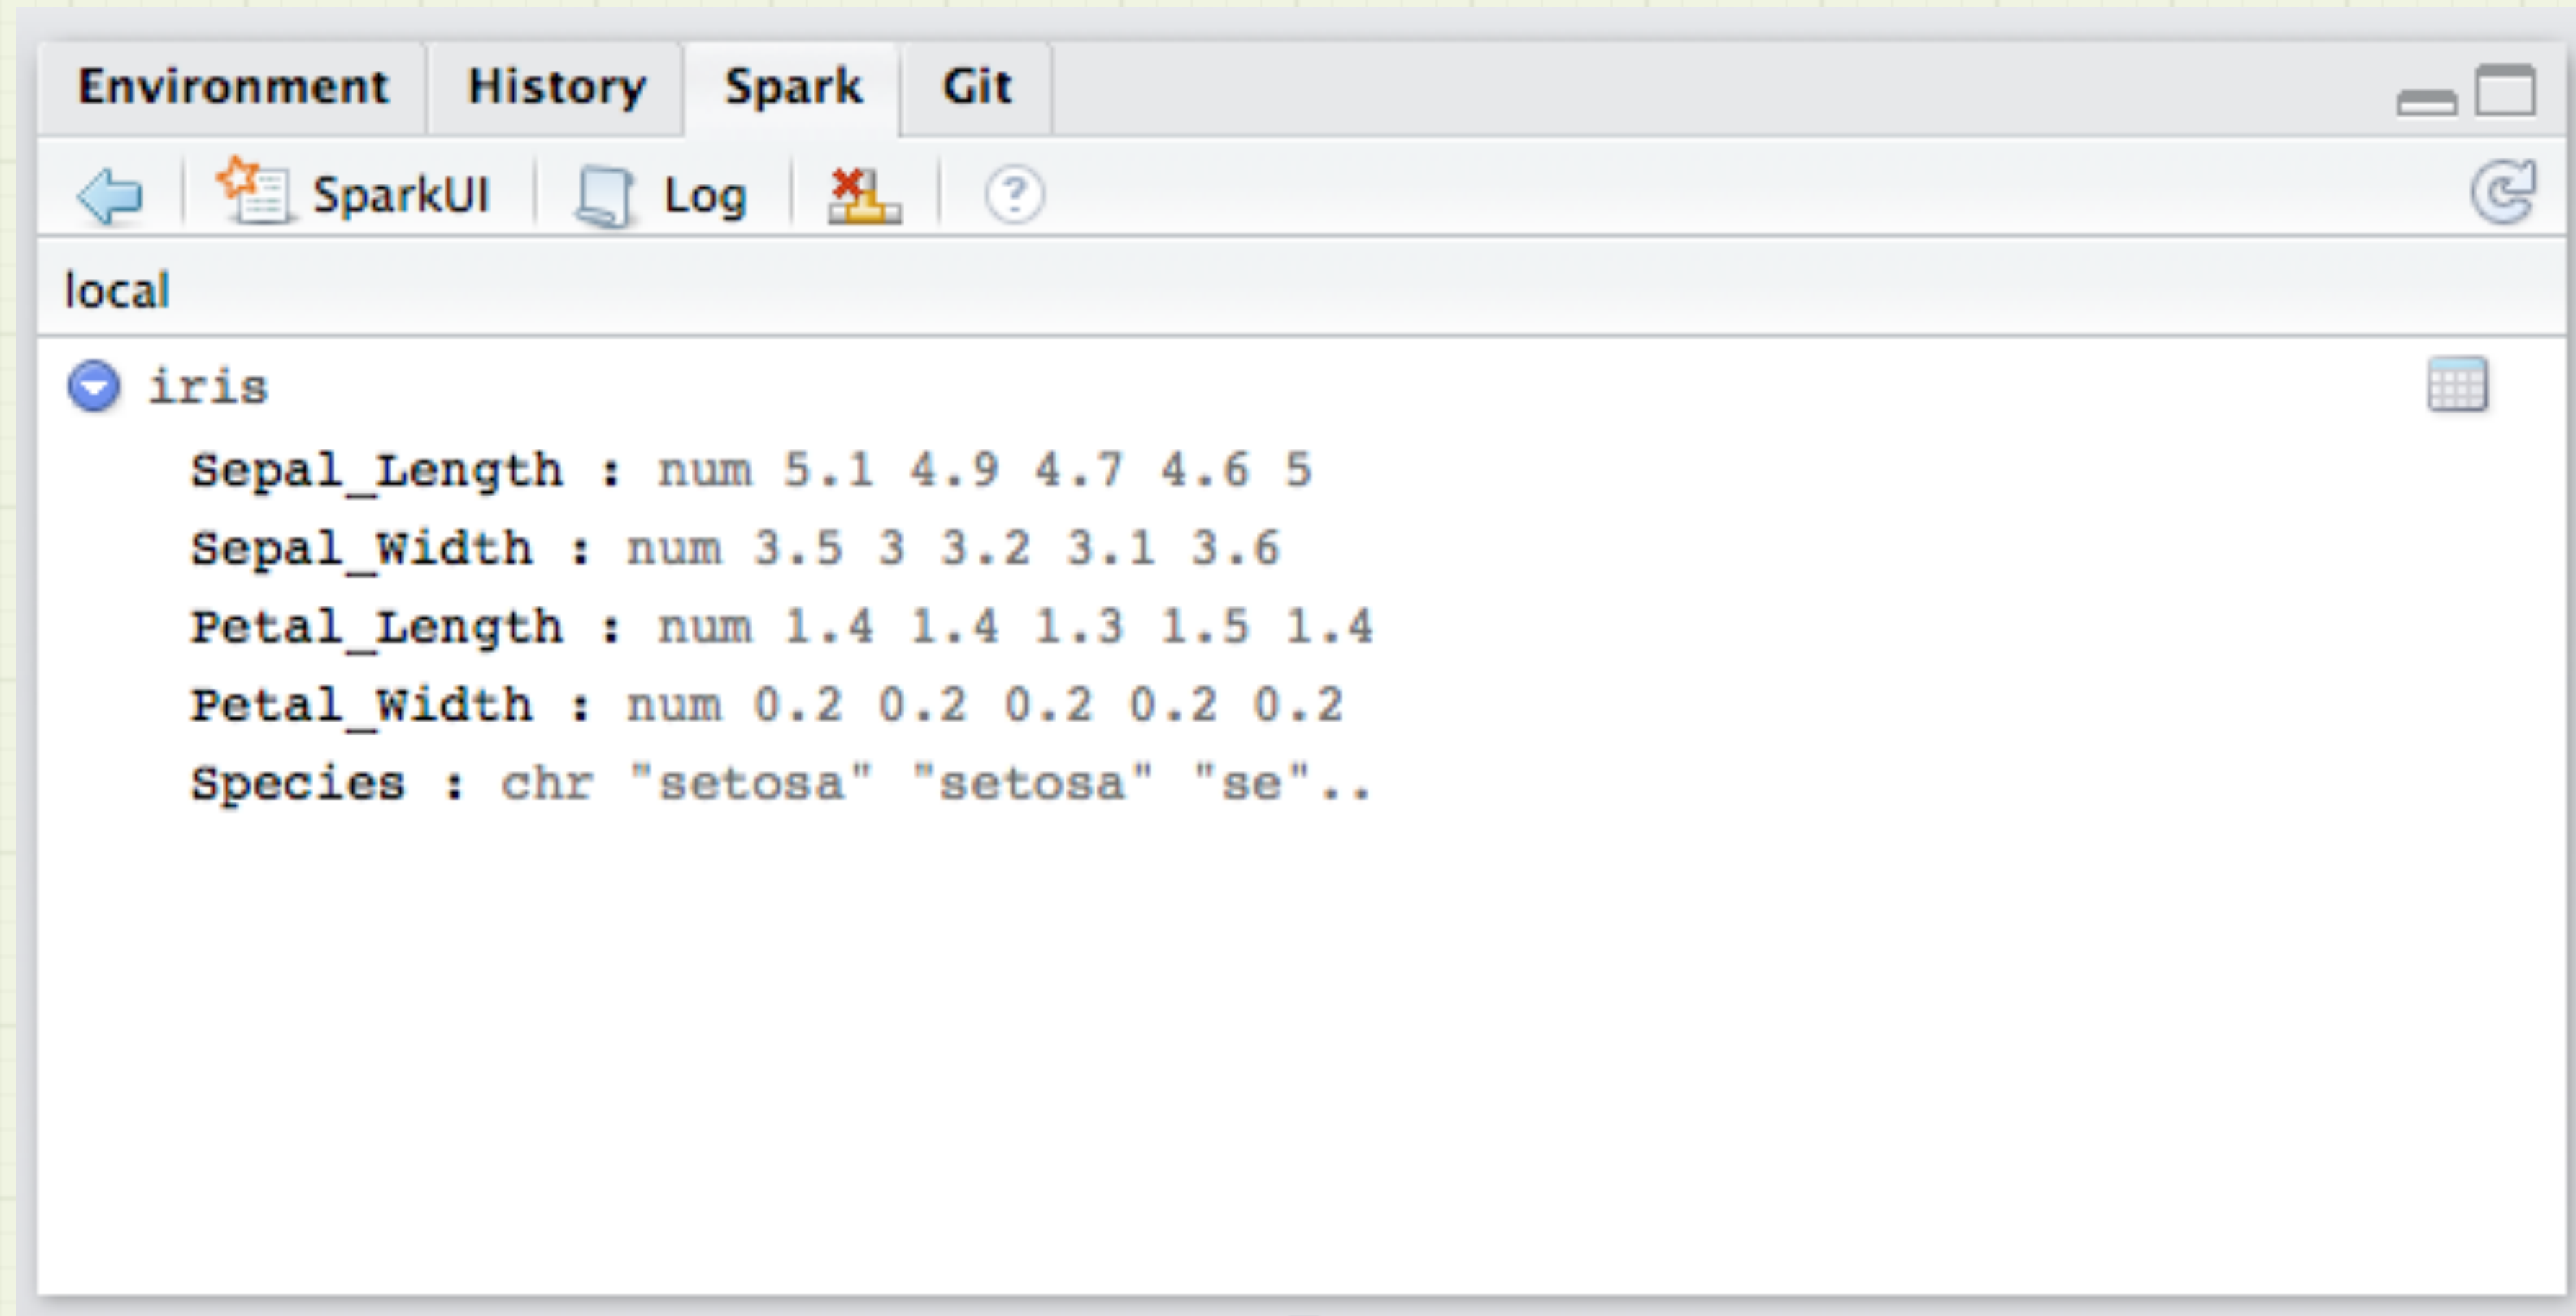
	Length	Class	Mode
src 1		src_spark	list
ops 3		op_base_remote	list

```
> broom::glance(iris_tbl)
```

```
Error: glance doesn't know how to deal with data  
of class tbl_sparktbl_sqltbl_lazytbl
```



fix: RStudio Spark browser



optional (experimental/expensive) fix: `replyr::replyr_summary()`

```
> library("replyr")  
> replyr_summary(iris_tbl, countUniqueNonNum= TRUE)
```

	column	index	class	nrows	nna	nunique	min	max	mean	sd	lexmin	lexmax
1	Sepal_Length	1	numeric	150	0	NA	4.3	7.9	5.843333	0.8280661	<NA>	<NA>
2	Sepal_Width	2	numeric	150	0	NA	2.0	4.4	3.057333	0.4358663	<NA>	<NA>
3	Petal_Length	3	numeric	150	0	NA	1.0	6.9	3.758000	1.7652982	<NA>	<NA>
4	Petal_Width	4	numeric	150	0	NA	0.1	2.5	1.199333	0.7622377	<NA>	<NA>
5	Species	5	character	150	0	3	NA	NA	NA	NA	setosa	virginica



Remote data is not “in the tidyverse”

```
> library("tidyr")  
> iris %>% nest(-Species)
```

```
# A tibble: 3 × 2
```

Species	data
<fctr>	<list>
1 setosa	<tibble [50 × 4]>
2 versicolor	<tibble [50 × 4]>
3 virginica	<tibble [50 × 4]>

```
> iris_tbl %>% nest(-Species)
```

```
Error in UseMethod("nest_") :  
no applicable method for 'nest_' applied to an object  
of class "c('tbl_spark', 'tbl_sql', 'tbl_lazy', 'tbl')"
```



To use the “The Split-Apply-Combine Strategy for Data Analysis”

- You must use `dplyr::group_by()`
- And restrict yourself to operators that are “group aware.”

```
> iris_tbl %>%  
  group_by(Species) %>%  
  summarize_all(funs(typical=mean))
```

Source: query [3 x 5]

Database: spark connection master=local[4] app=sparklyr local=TRUE

	Species	Sepal_Length_typical	Sepal_Width_typical	Petal_Length_typical	Petal_Width_typical
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	versicolor	5.936	2.770	4.260	1.326
2	virginica	6.588	2.974	5.552	2.026
3	setosa	5.006	3.428	1.462	0.246

How is all this implemented?

- `dplyr` uses `sparklyr` to translate each `dplyr`-verb into SparkSQL.
- `dplyr` collects a compound query representing arbitrarily long sequences of operations on the remote (Spark) data.
- The query is lazy, and only run if and when the data is actually used in an eager calculation (such as printing).

Query examples

```
> iris_tbl %>%  
  group_by(Species) %>%  
  summarize_all(funs(typical=mean)) %>%  
  show_query
```

```
<SQL>  
SELECT `Species`, AVG(`Sepal_Length`) AS `Sepal_Length_typical`, AVG(`Sepal_Width`) AS  
`Sepal_Width_typical`, AVG(`Petal_Length`) AS `Petal_Length_typical`, AVG(`Petal_Width`) AS  
`Petal_Width_typical`  
FROM `iris`  
GROUP BY `Species`
```

```
> iris_tbl %>%  
  head %>%  
  show_query
```

```
<SQL>  
SELECT *  
FROM `iris`  
LIMIT 6
```

Being eager

- `dplyr::collapse()`
 - Try and simplify the the accumulated query.
- `dplyr::compute()`
 - Force computation, materialize result of query into a new table.
- `dplyr::collect()`
 - Force computation, materialize result of query into a local `tbl`.

Lazy eval: nothing done until something triggers compute()

```
res <- iris_tbl %>%  
  group_by(zspecies) %>%      #OOOPS!  
  summarize_all(funs(typical=mean))  
# no error!
```

```
print(res)
```

Error:

org.apache.spark.sql.AnalysisException:
cannot resolve '`zspecies`' given

```
res <- iris_tbl %>%  
  group_by(zspecies) %>%  
  summarize_all(funs(typical=mean)) %>%  
  compute()
```

Error:

org.apache.spark.sql.AnalysisException:
cannot resolve '`zspecies`' given

```
res <- iris_tbl %>%  
  group_by(Species) %>%  
  summarize_all(funs(typical=mean)) %>%  
  show_query
```

<SQL>

SELECT `Species`, AVG(`Sepal_Length`) ...

```
res <- iris_tbl %>%  
  group_by(Species) %>%  
  summarize_all(funs(typical=mean)) %>%  
  compute(name= 'summarizedIris') %>%  
  show_query
```

<SQL>

SELECT *
FROM `summarizedIris`



Be cautious with remote data

- Semantics may be different than R conventions.
- May have limited ability to represent NA/NULL.
- Will not be able to represent factors other than as strings.
- May have different integer and floating point arithmetic rounding and rules.
- May have different column names and quoting/escaping conventions.
- Do not accept arbitrary R user define functions.
 - Will substitute many common R functions by name.
- May not have easy to access ranking and other window functions.



The biggest difference

- Most remote data sources do not:
 - Guarantee row order
 - Support row-names
- These are deliberately *not* included as relational concepts.
- R local data frames guarantee row order
 - Many calculations depend on this
 - These calculations will not be correct on remote data sources.



dplyr::mutate example

```
> iris_tbl %>%  
  summarize(mx = mean(Sepal_Length))
```

```
Source:   query [1 x 1]
```

```
Database: spark connection master=local[4] app=sparklyr local=TRUE
```

```
      mx  
    <dbl>  
1 5.843333
```

```
> iris %>%  
  summarize(mx = median(Sepal.Length))
```

```
      mx  
1 5.8
```

```
> iris_tbl %>%  
  summarize(mx = median(Sepal_Length))
```

```
Error: org.apache.spark.sql.AnalysisException: Undefined function: 'MEDIAN'. This func  
neither a registered temporary function nor a permanent function registered in the da  
'default'.; line 1 pos 7
```



Spark union example

```
> d1 <- data_frame(x= 1, y= 'a')
> d2 <- data_frame(y= 'b', x= 2)
```

```
> bind_rows(d1, d2)
```

```
# A tibble: 2 × 2
```

	x	y
	<dbl>	<chr>
1	1	a
2	2	b

```
> union(d1, d2)
```

```
# A tibble: 2 × 2
```

	x	y
	<dbl>	<chr>
1	2	b
2	1	a

```
> d1s <- copy_to(sc, d1)
```

```
> d2s <- copy_to(sc, d2)
```

```
> bind_rows(d1s, d2s)
```

```
Error: incompatible sizes (1 != 2)
```

```
> union(d1s, d2s)
```

```
Source:   query [2 x 2]
```

```
Database: spark connection
```

```
master=local[4] app=sparklyr local=TRUE
```

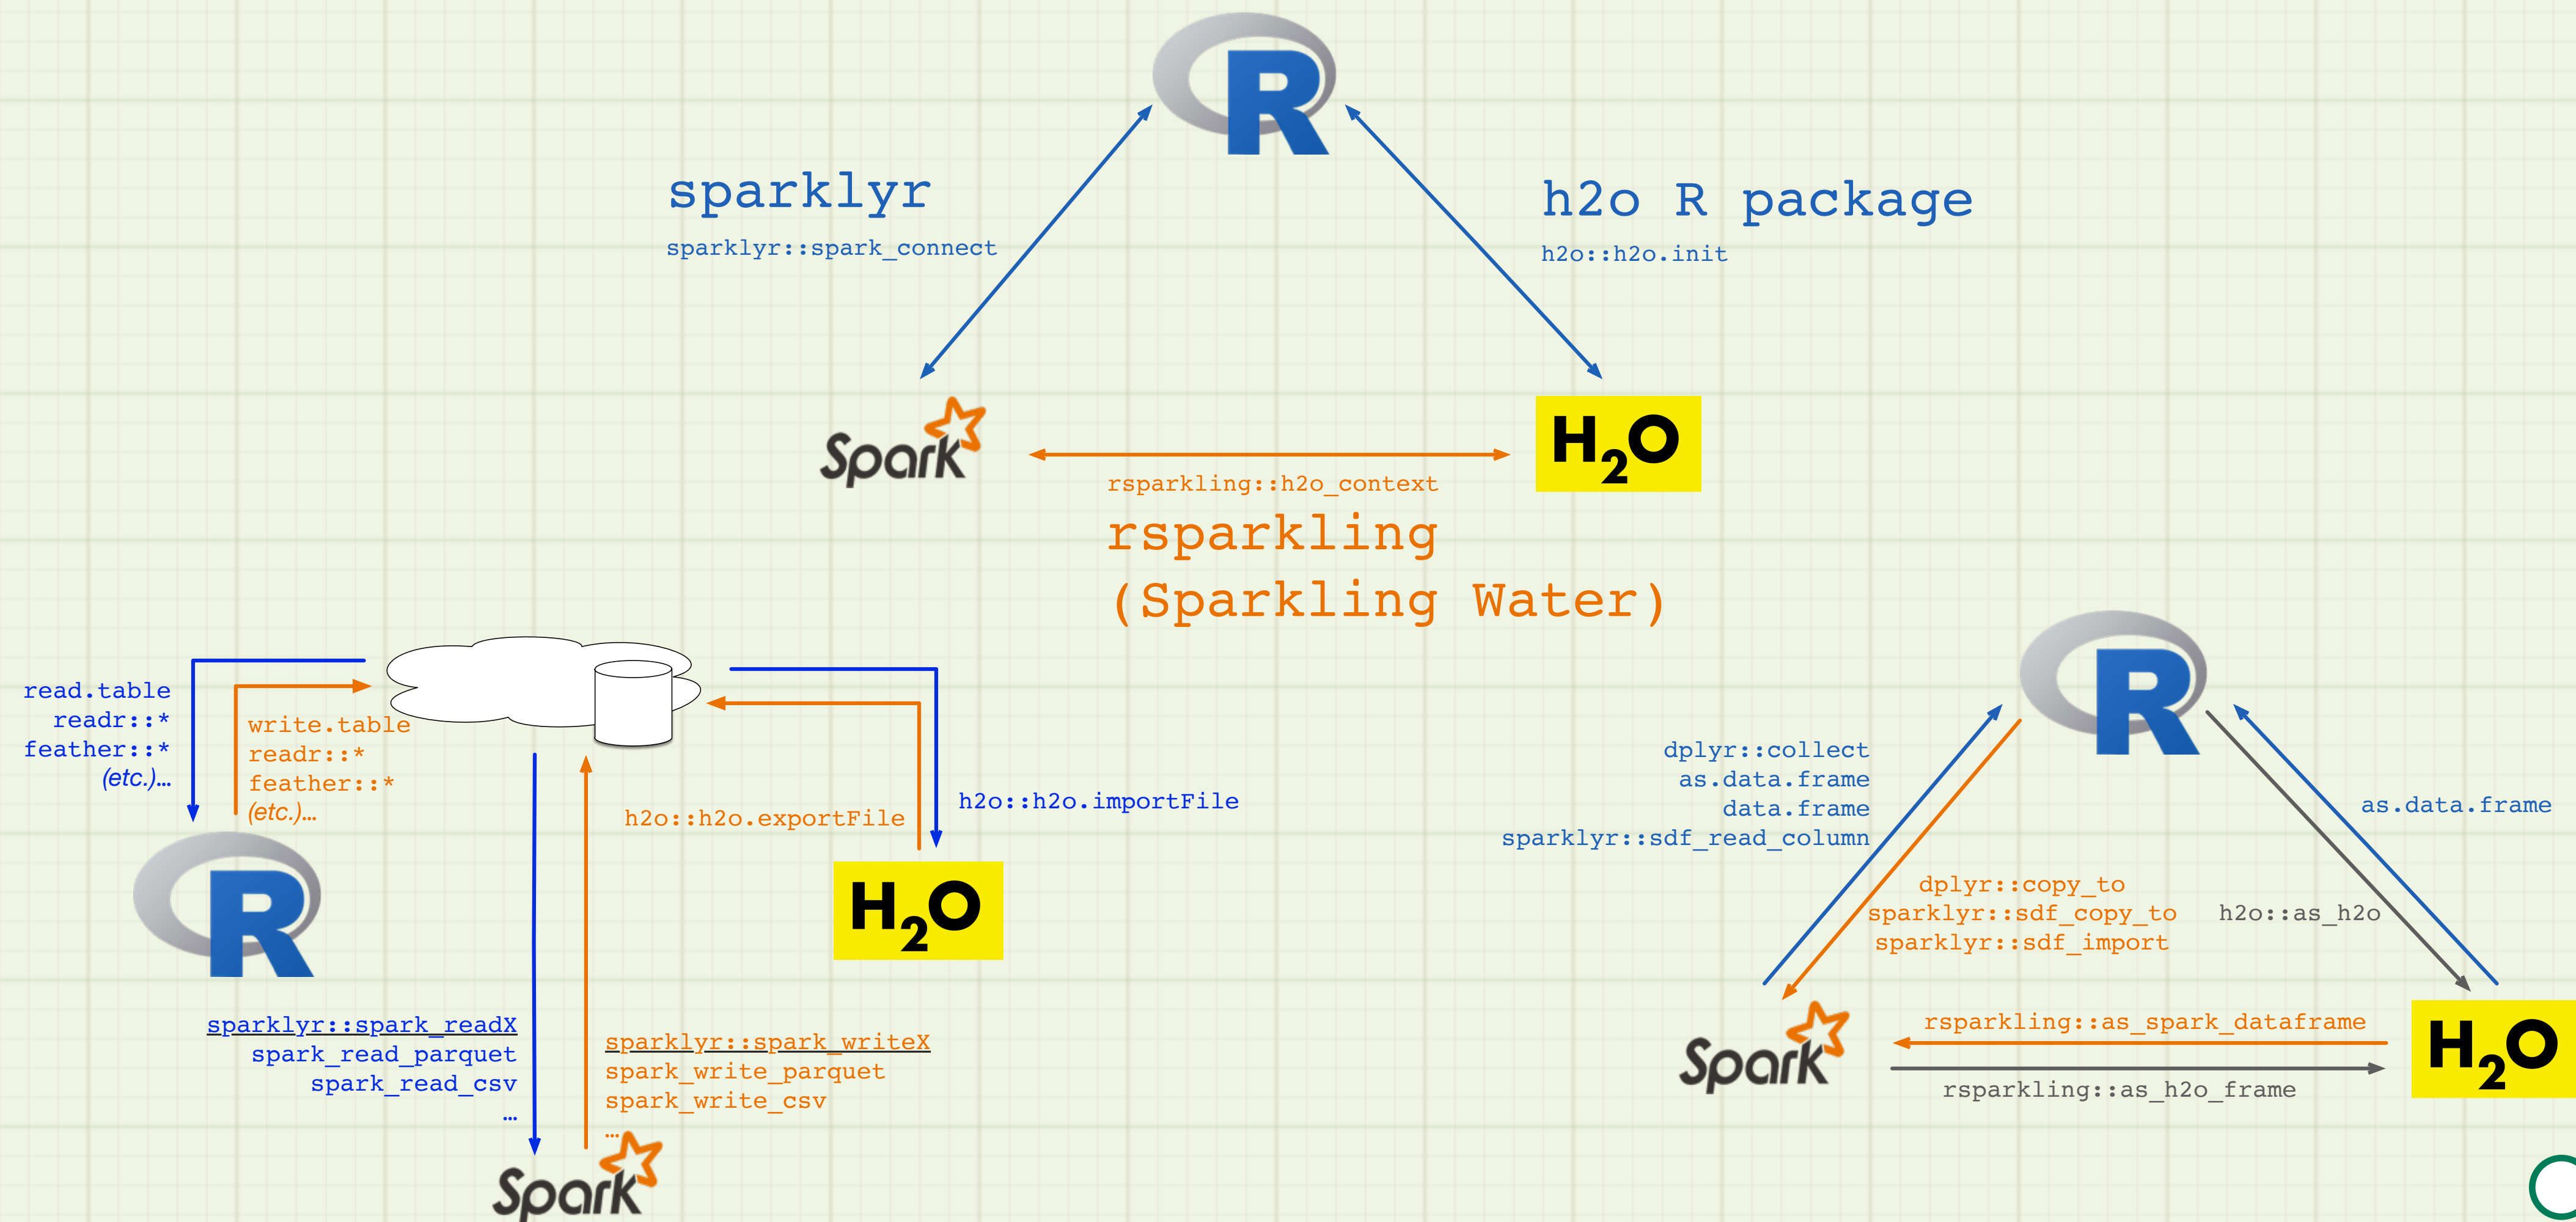
	x	y
	<chr>	<chr>
1	b	2.0
2	1.0	a



Some commands to remember

Command	What it does
<code>DBI::dbListTables(sc)</code>	Lists data items in Spark
<code>dplyr::tbl(sc, 'iris')</code>	Build a handle pointing to the Spark object with the given name
<code>dplyr::copy_to(sc, iris, "iris")</code>	Copy data from R to Spark and choose name of result
<code>dplyr::collect(iris_tbl)</code>	Copy data from Spark to R
<code>dplyr::db_drop_table(sc, 'iris')</code>	Drop object by name

Please keep the “three island slides” as a handy reference



Quality of implementation varies by service provider

Service provider	R package	Quality of the experience
SQLite	RSQLite	excellent
PostgreSQL	RPostgreSQL	excellent
Spark2.x.x	SparklyR	excellent
Spark1.6.2	SparklyR	passable
MySQL	RMySQL	low

Also consider using Spark-specific commands

- **sdf_*** : Use spark data frame commands to manipulate data frames.
 - Examples: sdf_partition, sdf_predict, sdf_sample
- **ft_*** : Use feature transforms to manipulate features.
 - Examples: ft_bucketizer, ft_index_to_string
- **ml_*** : Use machine learning algorithms to train models.
 - Examples: ml_kmeans, ml_logistic_regression, ml_pca

Break: please be back by 3:30 pm.

Then we work through some Spark
exercises together